VRIJE
UNIVERSITEIT
BRUSSEL

DISTRIBUTING REACTIVE HAAI

Haai Virtual machine

Hans

June 2024

Promotors: prof. dr. Wolfgang De Meuter   Supervisor: Bjarno Oeyen

sciences and bioengineering sciences

ii

Abstract

We are looking into the distribution of a purely reactive programming language Haai.
Context . **inquiry** . **Approach** . **Knowledge** . **Grounding** . **Importance**

Contents

Chapter 1

Introduction

This thesis steps into the world of distributed reactive programs. Reactive programming revolves around the concept of reacting on streams of data. In reactive programming the program reacts on a stream of data. The stream of data travels trough the reactive program to exit in a possibly altered state. We will investigate possibilities to distribute the reactive programming language Haai. (Oeyen et al., 2024) Haai is a purely reactive language consisting of only reactors. The literature review on distribution models for reactive programs lead us to the actor model (De Koster et al., 2016). We have implemented a virtual machine that opens the possibilities for the Haai language to be distributed on a Elixir (Juric, 2024) cluster.

1.0.1 Statement of the problem

To distribute a reactive program different approaches can be taken, We will select a distribution model to experiment with the distribution of reactive programming languages or reactors written in Haai. Glitch avoidance in reactive programs has been achieved on most modern implementations that run on a single node. In a distributed setting glitch avoidance is not an easy task due to the parts the distributed system is build of. Delays, network failures and other problems make it complicated to reach glitch freedom in a distributed setting. (Bainomugisha et al., 2013)

1.0.2 Definition of terms

Reactive programming language A computer language to write reactive programs.

Reactive programming Reactive programming is a programming paradigm that focuses on asynchronous data streams and the propagation of changes. In reactive programming, you model your application as a series of data flows and transformations, where changes in one part of the system can automatically trigger updates in other parts. This paradigm is particularly useful for building interactive and responsive applications, such as user interfaces, real-time analytics, and event-driven systems.

Distribution the deployment of reactive systems across multiple computing nodes or devices. This distribution enables the system to handle a large volume of data, support high concurrency, and scale horizontally to meet increasing demands.

1.0.3   Remaining chapters

In the next chapter we give a description of the purely reactive programming language Haai. Chapter 3 about the main implementation of this thesis describes the Haai virtual machine (Hvm). The distribution chapter explains why we selected Elixir as a platform and how we build the distributed cluster that runs multiple independent Hvm's. Before concluding and reflecting on further work we have a description of the use case of the thesis.

Chapter 2

Haai

We are making use of the novel purely reactive language Haai. Haai is reactive all-the-way-trough. As such the only construct available for expressing computations is a reactor. A reactor has sources or inputs and sinks or outputs. For each time the sources are updated the reactor will react and produce a new sink value. The number of sources and sinks are variable per reactor definition. In Haai there are no functions only reactors. This allows us to build reactive programs without functions. (Oeyen et al., 2024)

Two common problems encountered within reactive programming are by design eliminated. These two problems aer called the *Reactive Thread Hijacking Problem* and the *Reactive/Imperative Impedance Mismatch.*(Vonder et al., 2020) Two problems that arise from mixing reactive code with sequential code in a two-layered reactive programming language.

Reactive Thread Hijacking Problem

This happens in a two-layered reactive programming language when the reactive runtime or top layer calls a function in the bottom layer and that call never returns. The host has hijacked the control and the reactive runtime will not be able to regain control.

Reactive/Imperative Impedance Mismatch

When the reactive runtime in a two-layered reactive program used different side effects by calling functions on the first layer. It might be so that the order in which they are executed is not as expected. Creating faults in the programs that might not be easy to detect.

2.0.1   compare with other rx

2.1   Haai instruction set

After compiling a Haai program into its bytecode everything is based around this instruction set. The instruction set defines the set of low-level operations that the virtual machine can execute. Each instruction corresponds to a specific operation. the Haai virtual machine presented in the next chapter is build around this instruction set.

[**I-ALLOCMONO, reactor**] Allocate memory for the given reactor.

[**I-LOOKUP, signal**] Lookup the value for signal in a signal table.

[**I-SUPPLY, from, destination, index**] Supply the value found in from to destination at
     index.

[**I-REACT, memory_location**] Apply the reaction required at memory_location.

[**I-CONSUME, memory_location, index**] Move the value form REACT to the run-time-
     memory.

[**I-SINK, rti_location, sink_index**] Remove value from run-time-memory and store in the
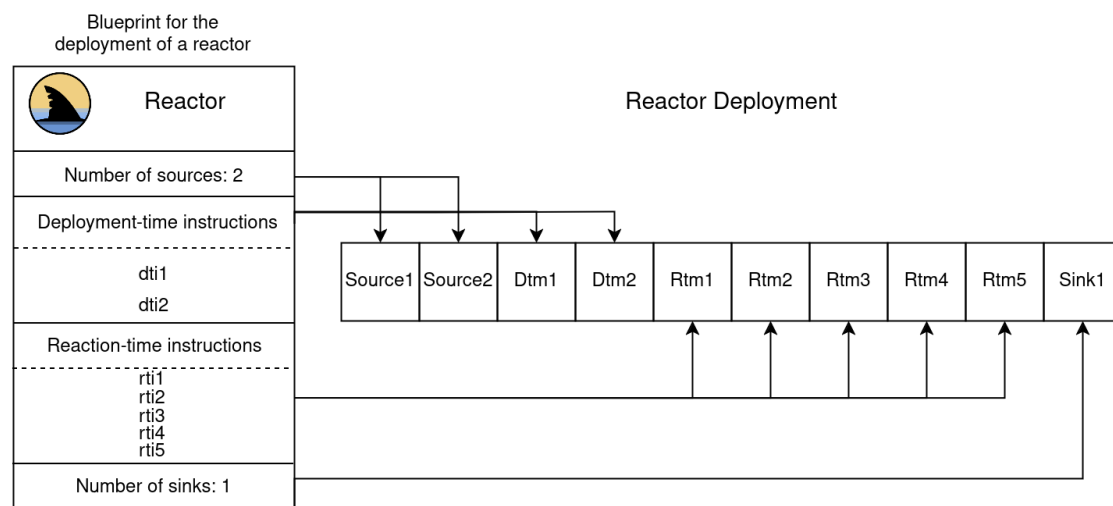   reactor sink.

2.2   Reactors

In Haai, a purely reactive computer language, the only construct available for expressing com-
putations is a reactor. We can write one simple reactor or combine multiple reactors in one
'main' reactor or program. The main program is build out of reactors, these can be native or
user defined. There is no underlying layer from which we would lift functions and so the Haai
virtual machine provides for example native reactors for basic arithmetic. The deployment of a
reactor reacts on signals received, that signal altered by the reaction is found in the sink of the
reactor. The deployment of a reactor demand for memory block fixed in size. The size of the
total memory required for the deployment of a reactor is mainly influenced by the number of run
time instructions the reactor has.

2.2.1   Deploying a reactor

A reactor can be seen as a blue print for its deployment, similar how an object created from
a class in object-oriented programming is an instantiation of the blueprint the class represents.
The deployment of the reactor will request a certain size of memory. That total block of memory
exists out of four parts: Sources, deployment time memory (dtm), reaction time memory (rtm)
and the sinks. When we go from a reactor to a deployment we can find the following information
in the reactor or blue print for the deployment. The number of sources and sinks, the deployment
time instructions (dti) and the reaction time instructions (rti). For each source and each sink
a position is reserved in the memory block we are constructing. From the deployment time
instructions in the reactor we can prepare the deployment time memory (dtm). The Dtm part
of the memoryblock we are constructing will reserve space for the deployment of the reactors
found in the dti.

   Now we have reserved space for the sources the sinks and dtm, the last step is to reserve
space for the reaction time memory (rtm). The space required for rtm is equal to the number of
instructions found in the reaction time instructions. For each instruction in rti there is one block
reserved in rtm. A deployed reactor uses the four memory parts: sources, dtm, rtm and sinks to
react according to the reaction time instructions found in rti. The following figure expresses the
deployment of a reactor.

Blueprint for the
deployment of a reactor

Reactor

Number of sources: 2

Deployment-time instructions

dti1
dti2

Reaction-time instructions

rti1
rti2
rti3
rti4
rti5

Number of sinks: 1

Reactor Deployment

Source1 | Source2 | Dtm1 | Dtm2 | Rtm1 | Rtm2 | Rtm3 | Rtm4 | Rtm5 | Sink1

Chapter 3

Haai virtual machine

To make it possible for Haai code to work on multiple platforms one can make a virtual machine for each required platform. After the example of Remus, a virtual machine that has been carefully designed to be usable for running reactive programs in low-powered computing environments.(Oeyen et al., 2022b) We implement in this thesis a Haai virtual machine to experiment with the distribution of Haai. The Haai virtual machine (Hvm) is developed in the Elixir computer language. The chapter on distribution arguments our choice for Elixir.

3.1 Running the reactor

The Haai virtual machine needs the following arguments to start:

reactor bytecode Bytecode obtained from the compilation of the reactor with the Haai compiler.

connect source x A function that the Hvm can use to connect input to the sources of the reactor. At the start of the next iteration this function will be called to update the sources of the deployed reactor.

handle sink the function that after each iteration is called to handle the new sinks rproduced by the deployed reactor.

3.2 Managing the reactors state

3.3 Example bytecode program

After parsing the bytecode that was compiled form the original Haai code with the Haai compiler into an Elixir readable format of nested lists a program could look like this. In this example we see two user defined reactors named consonance and note_length followed by the main reactor that functions as the main reactor to be executed.

```
co_nl = [
[:consonance, 1, 1,
[
    ["I-ALLOCMONO", :multiply]
```

```
],
[
    ["I-LOOKUP", :ci],
    ["I-SUPPLY",["%RREF",1],["%DREF",1],1],
    ["I-SUPPLY",["%SRC",1],["%DREF",1],2],
    ["I-REACT",["%DREF",1]],
    ["I-CONSUME",["%DREF",1],1],
    ["I-SINK",["%RREF",5],1]]
],

[:note_length, 1, 1,
[
    ["I-ALLOCMONO", :divide],["I-ALLOCMONO", :multiply]
],
[
    ["I-SUPPLY",60000,["%DREF",1],1],
    ["I-SUPPLY",["%SRC",1],["%DREF",1],2],
    ["I-REACT",["%DREF",1]],
    ["I-LOOKUP",:lm],
    ["I-SUPPLY",["%RREF",4],["%DREF",2],1],
    ["I-CONSUME",["%DREF",1],1],
    ["I-SUPPLY",["%RREF",6],["%DREF",2],2],
    ["I-REACT",["%DREF",2]],
    ["I-CONSUME",["%DREF",2],1],
    ["I-SINK",["%RREF",9],1]]
],

[:main, 2, 2,
[
    ["I-ALLOCMONO", :consonance],["I-ALLOCMONO", :note_length]
],
[
    ["I-SUPPLY",["%SRC",1],["%DREF",1],1],
    ["I-REACT",["%DREF",1]],
    ["I-SUPPLY",["%SRC",2],["%DREF",2],1],
    ["I-REACT",["%DREF",2]],
    ["I-CONSUME",["%DREF",1],1],
    ["I-SINK",["%RREF",5],1],
    ["I-CONSUME",["%DREF",2],1],
    ["I-SINK",["%RREF",7],2]
]
]
]
```

The above is an example of bytecode that is send to the Haai virtual machine. This bytecode is a compilation of the following Haai program defining two user defined reactors and joining them together within the main reactor. The main reactor could be seen as the 'main' function in C. The starting point or the call that kicks off the programs execution.

(

```
(defr (consonance f)
     (out (* ci f)))

   (defr (note_length bpm)
      (def q (/ 60000 bpm))
          (out (* lm q)))

   (defr (main f bpm)
      (out (consonance f))
      (out (note_length bpm)))
)
```

3.4 Native reactors

The Haai virtual machine makes use of native reactors to perform simple arithmetic operations like addition, division, multiplication and subtraction. Native reactors are recognized by the Hvm and applied when required. No lifting, so we need these reactors cannot perform side effects.. to perform IO there are special reactors outside of the program (2.5 in witouth funcitons)

Chapter 4

Distribution

4.1 Shared-nothing architecture

To organize the distribution of the Haai virtual machine we have opted for a shared-nothing architecture. (Stonebraker, 1986) This allows us to build the cluster with the following characteristics, suitable for the first time distribution of the reactive language Haai.

Independent Nodes Each node can execute its own specific code and perform computations independently since no memory, storage or processing power is shared.

No shared state Each node manages its state independently. Cluster communication between nodes is done with messages. we use Erlang distributed mechanisms in the implementation.

Horizontal scalability Shared-nothing architectures are highly scalable horizontally, this makes it straightforward to add or remove nodes from the cluster. It allows us to run clusters of all sizes. Practically the cluster size will be bound to the available hardware.

Fault isolation Faults or failures in one node do not affect the operation of other nodes in the system. We can monitor node failure with Erlang distributed mechanisms and act accordingly.

Flexibility and Modularity Changes to one node do not impact the operation of other nodes, allowing for maintenance, upgrades, and system evolution of the cluster.'

4.2 Erlang virtual machine

We implemented the shared-nothing architecture in Elixir a functional, concurrent, and fault-tolerant programming language built on the Erlang Virtual Machine (BEAM). The Erlang Virtual Machine provides a robust platform due to several key features and design principles:

Concurrency model Erlang has a lightweight concurrency model which is based on actor-like processes. Since each process runs independently and can communicate with other processes trough message passing the Erlang virtual machine is inherently decentralized by design.

Fault tolerance The BEAM VM provides mechanisms for process supervision, fault isolation, and hot code swapping, allowing systems to recover from failures without affecting the overall operation of the cluster.

Distribution transparency Erlang's distribution features enable processes to communicate
effortlessly across network boundaries. Within an Erlang cluster, nodes can seamlessly
interact through remote procedure calls (RPCs) and message passing. Different types of
distribution models can be build on the Erlang platform.

Scalability Horizontal scaling or adding nodes to a cluster is straightforward. This makes it
easy to scale clusters to need.

Tooling Common tasks are simplified thanks to a comprehensive tool set. tools like the Mix
build tool and the OTP framework are core tools used in the implementation of the dis-
tributed Haai cluster.

The example program is explained in more detail in the use case chapter. The same program
or reactor can be run many times. The distributed setup we make exists out of a number of elixir
nodes that all send there messages to one sound server. The bigger the distributed network is
the more messages that are send to the sound server. The sound server it self has virtually no
limit to receive messages. The main parts of the distribution are taken care of by the underlying
Beam virtual machine. One main reason to develop the Hvm on this platform is to be able to
distribute in a solid and proven way.

4.3   Actor model

The actor system in Erlang implements the processes actor model. (De Koster et al., 2016)
These actors are modeled as processes and can define there receive primitive to specify what
messages it responds to. Thanks to the isolated turn principle the actor model guarantees to be
free of data races and deadlocks by design. In Erlang actors are typically used when you need to
encapsulate state and behavior within a single unit of computation., actors are spawned using
constructs like GenServer, Task, or Agent within a single node. All actors communicate with
each other using message passing within the same node. One can use these actors for modeling
concurrent and stateful components in an application.

## 4.4   Distributed Erlang mechanisms

Distributed Erlang mechanisms come into play when you need to communicate between actors
residing on different nodes within a cluster. Nodes communicate with each other using distributed
Erlang protocols, allowing actors on different nodes to exchange messages. Distributed Erlang
mechanisms provide fault tolerance, distribution transparency, and location transparency across
the cluster.

Distributed Erlang provides location transparency, allowing actors to communicate regardless
of their physical location within the cluster. This transparency hides the underlying network
details, promoting loose coupling (Carreton et al., 2010) between nodes and actors. Actors on
different nodes can interact seamlessly through message passing, without needing to be aware of
the specific nodes they are communicating with. Loose coupling in distributed Erlang enables
the system to scale horizontally by adding or removing nodes without significantly impacting the
overall architecture.

Chapter 5

Distributed reactive melody generator

The use case of this thesis makes sound. The distributed Haai cluster, existing out of many distributed reactors provides the sink values form the reactors as values inserted into OSC messages (Schmeder et al., 2010) that are send to a sound server to trigger the synthesis of some sound. It is the reactors that, by reacting on incoming streams of data, produce the values that result in a rhythmical flow of OSC messages to the sound server.

The flow of rhythmical messages is related to the duration of each sound. That same duration that makes the sound disappear after it was started determines the iteration speed for the reactor or program that runs on each individual node. Each reactor iterates over and over but each iteration has a defined duration, this duration is virtually equal to the duration for a sound calculated by the reactor. A reactor will calculate the new values, sleep for the calculated duration and then provide the new values in the sink. At this point the Hvm calls a function that it was given to handle the sink values. For this use case that function constructs the message and sends it to the sound server that propagates it to the requested synthesizer.

5.0.1   Sound server

The synthesizer that produces the sound is a Synthdef (the definition of a synthesizer) inside the sound server named SuperCollider. Supercollider is an open-source platform for audio synthesis and algorithmic composition. Developed by James McCartney in the late 1990s (Wilson et al., 2011), it has since become a powerful tool in the fields of music technology, computer music, and sound art. Supercollider provides a flexible and expressive environment for creating and manipulating sound in real-time, making it an invaluable resource for both research and artistic exploration. Some main features of super collider include:

Audio Synthesis Supercollider offers a wide range of synthesis techniques, including additive, subtractive, granular, and physical modeling synthesis. Users can create complex sounds by combining these techniques and modulating parameters in real-time.

Real-time Processing One of Supercollider's key strengths is its ability to process audio in real-time. This makes it suitable for live performances, interactive installations, and other time-sensitive applications.

Algorithmic Composition Supercollider provides tools for generating music algorithmically, allowing users to create compositions based on mathematical algorithms, rulesets, or generative processes.

Integration with External Hardware Supercollider can interface with external MIDI controllers, audio interfaces, and other hardware devices, enabling users to incorporate physical instruments and sensors into their sound projects.

Community and Documentation Supercollider has a vibrant online community of users who share code, tutorials, and resources. Additionally, comprehensive documentation is available, including tutorials, reference guides, and examples to help users learn and master the software.

SuperCollider has been and still is actively used in academia for many purposes, for example:

Research in Music Technology Digital signal processing, human-computer interaction, and machine learning for music. Researchers leverage its flexibility and programmability to prototype new algorithms, experiment with novel synthesis techniques, and investigate the perceptual and cognitive aspects of sound.

Composition and Sound Design Supercollider is used to create innovative works that push the boundaries of traditional music and sound art. Its ability to generate complex and evolving textures, as well as its support for algorithmic composition, makes it a valuable tool for exploring new sonic territories.

Teaching and Learning Supercollider is increasingly incorporated into music technology and computer music curricula at universities and colleges worldwide. It provides students with hands-on experience in sound synthesis, programming, and digital audio processing.

5.0.2 Sound server communication

The communication between the nodes in the Elixir cluster and the sound server is done over the network using the UDP network protocol. The sound server can handle a virtual unlimited amount of messages, but is practically bounded to the udp buffer size on the actual machine that runs the sound server. As such the Elixir cluster of reactors can be of any feasible size and the sound server will respond on any number of messages by synthesizing the requested sound.

5.0.3 Musical values

The reactor or program that runs on all nodes in the cluster does for each note the following. It calculates two values given two input streams of numbers. The first source is a stream of base frequencies, represented as a float number, the second source is again a stream of numbers representing a tempo in beats per minute (bpm)

Consonant notes

Consonant notes sound harmonious when played together. consonance can be expressed as a ratio between two notes or frequencies. Ratios like the perfect fifth (frequency ratio of 3:2) and the major third (frequency ratio of 5:4) are very commonly used ratios in a musical setting. One of the two reactors in the main program used for this use case calculates such a note for a given base note, both are expressed in there actual frequency represented as a float number.

Tempo

We can express the speed of a musical note progression in beats per minute (bpm). From that number we can then calculate the duration of individual notes in the musical performance.

We calculate the time in milliseconds. Since one minute is 60000 milliseconds one can calculate the length of a quarter note with $q = \frac{60000}{bpm}$ with the assumption that we use a 4/4 time signature. From the length of the quarter note it is easy to find other durations. For example, the duration of a halve note is double that of a quarter note or the duration of a eight note is half the time of a quarter note.

One of the two reactors in the main program used for this use case calculates exactly that, the length for a quarter note and improvises a deviser or multiplier to produce some duration.

Further work

Try other distribution models

Conclusion

Bibliography

Bainomugisha, E., Carreton, A. L., Cutsem, T. V., Mostinckx, S., & Meuter, W. D. (2013). A survey on reactive programming. *ACM Comput. Surv.*, *45*(4), 52:1–52:34. https://doi.org/10.1145/2501654.2501666

Carreton, A. L., Mostinckx, S., Cutsem, T. V., & Meuter, W. D. (2010). Loosely-coupled distributed reactive programming in mobile ad hoc networks. In J. Vitek (Ed.), *Objects, models, components, patterns, 48th international conference, TOOLS 2010, málaga, spain, june 28 - july 2, 2010. proceedings* (pp. 41–60). Springer. https://doi.org/10.1007/978-3-642-13953-6_3

De Koster, J., Van Cutsem, T., & De Meuter, W. (2016). 43 years of actors: A taxonomy of actor models and their key properties. *Proceedings of the 6th International Workshop on Programming Based on Actors, Agents, and Decentralized Control*, 31–40. https://doi.org/10.1145/3001886.3001890

Juric, S. (2024). *Elixir in action*. Simon; Schuster.

Oeyen, B., De Koster, J., & De Meuter, W. (2022a). Reactive programming on the bare metal: A formal model for a low-level reactive virtual machine. *Proceedings of the 9th ACM SIGPLAN International Workshop on Reactive and Event-Based Languages and Systems*, 50–62.

Oeyen, B., De Koster, J., & De Meuter, W. (2022b). Reactive programming on the bare metal: A formal model for a low-level reactive virtual machine. *Proceedings of the 9th ACM SIGPLAN International Workshop on Reactive and Event-Based Languages and Systems*, 50–62. https://doi.org/10.1145/3563837.3568342

Oeyen, B., De Koster, J., & De Meuter, W. (2024). Reactive programming without functions. *The Art, Science, and Engineering of Programming*, *8*(3), 11. https://doi.org/10.22152/programming-journal.org/2024/8/11

Schmeder, A., Freed, A., & Wessel, D. (2010). Best practices for open sound control. *Linux Audio Conference*, *10*.

Stonebraker, M. (1986). The case for shared nothing. *IEEE Database Eng. Bull.*, *9*(1), 4–9. http://sites.computer.org/debull/86MAR-CD.pdf

Vonder, S. V. d., Renaux, T., Oeyen, B., Koster, J. D., & Meuter, W. D. (2020). Tackling the awkward squad for reactive programming: The actor-reactor model [ISSN: 1868-8969]. In R. Hirschfeld & T. Pape (Eds.), *34th european conference on object-oriented programming (ECOOP 2020)* (19:1–19:29). Schloss Dagstuhl–Leibniz-Zentrum für Informatik. https://doi.org/10.4230/LIPIcs.ECOOP.2020.19

Wilson, S., Cottle, D., & Collins, N. (2011). *The supercollider book*. The MIT Press.