



VRIJE
UNIVERSITEIT
BRUSSEL



DISTRIBUTING REACTIVE HAAI

Hans Van der Ougstraete

June 2024

Promotors: prof. dr. Wolfgang De Meuter Supervisor: Bjarno Oeyen

sciences and bioengineering sciences

Abstract

This thesis delves into the realm of distributed reactive programming, focusing on the distribution model for the purely reactive language Haai. Utilizing the actor model, the Haai language is distributed across a cluster of Elixir nodes, leveraging a shared-nothing approach to facilitate scalability and deployment. A validation of this distribution model is performed through an experiment involving a distributed reactive melody generator, which reacts to incoming streams of numbers by computing consonant frequencies and sound durations. This thesis addresses the challenge of distributed reactive programs, a task made complex by delays, network failures, and other distributed system issues. By exploring distribution models for reactive programming languages, this research contributes to understanding the intricacies of reactive systems in distributed settings.

Keywords: Reactive programming, Distribution, Haai, Elixir, Actor model, Distributed systems.

Contents

Abstract	iii
1 Introduction	1
1.1 Statement of the problem	1
1.2 Definition of terms	1
2 Reactive programming	3
2.1 Haai	4
2.2 Haai virtual machine	5
2.3 Remus instruction set	5
2.4 Reactors	5
2.5 Deploying a reactor	6
2.6 Running the reactor	7
2.7 Managing the reactors state	7
2.8 Native reactors	8
2.9 Special reactors	8
2.10 Example bytecode program	8
3 Distribution	11
3.1 Shared-nothing architecture	11
3.2 Erlang virtual machine	11
3.3 Actor model	12
3.4 Distributed Erlang mechanisms	12
3.5 Distributing the Hvm	13
4 Evaluation	15
4.1 Distributed reactive melody generator	15
4.2 Sound server	15
4.3 Sound server communication	17
4.4 Musical values	17
Conclusion	19

Chapter 1

Introduction

This thesis steps into the world of distributed reactive programs. Reactive programming revolves around the concept of reacting on streams of data. The language to distribute is Haai. (Oeyen et al., 2024) Haai is a purely reactive language consisting of only reactors. We chose for the actor model (De Koster et al., 2016) to distribute Haai code. The system runs on a cluster of Elixir (Juric, 2024) nodes with a shared-nothing approach. Allowing us to easily scale and deploy clusters of multiple sizes. We have validated the implementation with a sound generating experiment. The distributed reactive melody generator reacts on incoming streams of numbers and our reactors (programs in Haai) compute consonant frequencies and sound durations related to the input streams. This allowed us to evaluate the distribution model of the reactive programming language

1.1 Statement of the problem

To distribute a reactive program different approaches can be taken, We will select a distribution model to experiment with the distribution of reactive programming languages or reactors written in Haai. Glitch avoidance in reactive programs has been achieved on most modern implementations that run on a single node. In a distributed setting glitch avoidance is not an easy task due to the parts the distributed system is build of. Delays, network failures and other problems make it complicated to reach glitch freedom in a distributed setting. (Bainomugisha et al., 2013)

1.2 Definition of terms

Reactive programming language A computer language to write reactive programs.

Reactive programming Reactive programming is a programming paradigm that focuses on asynchronous data streams and the propagation of changes. In reactive programming, you model your application as a series of data flows and transformations, where changes in one part of the system can automatically trigger updates in other parts. This paradigm is particularly useful for building interactive and responsive applications, such as user interfaces, real-time analytics, and event-driven systems.

Distribution the deployment of reactive systems across multiple computing nodes or devices. This distribution enables the system to handle a large volume of data, support high concurrency, and scale horizontally to meet increasing demands.

In the next chapter we give a description of the purely reactive programming language Haai. Chapter 3 about the main implementation of this thesis describes the Haai virtual machine (Hvm). The distribution chapter explains why we selected Elixir as a platform and how we build the distributed cluster that runs multiple independent Hvm's. Before concluding and reflecting on further work we have a description of the use case of the thesis.

Chapter 2

Reactive programming

Reactive programming (RP) offers a paradigm that well fit the needs for building responsive, scalable and resilient software systems. RP offers a declarative approach for handling asynchronous streams of data. A reactive program will automatically react to changes over time. We can build (soft) real time applications that react given incoming data. We could say that the program has been programmed in a declarative way to know what to do when an input is presented to the ever running reactive program.

Some of the core principles of RP are:

Asynchrony: A key principle for reactivity is to not block the execution flow of the program. By handling events in an asynchronous way when they occur the flow of the program can continue and guarantee responsiveness.

Declarative composition: Composed declaratively a developer defines the desired behavior of the system without specifying the control flow of that system.

Data streams: RP treats data as a continuous stream, this enables the composition of complex data flows and allows the RP programs to react in real time to the continuous stream of data.

Backpressure handling: In RP we have producers and consumers for the data streams. It is critical to maintain a workable rhythm or flow between producers and consumers.

The above listed principles are used to guide the design and implementation of RP programs. Here under we highlight the key concepts that underpin the reactive programming paradigm.

Observable: An observable represents a source of asynchronous data streams that emit events over time. Observables can for example be created from network requests or timer events.

Observer: The observer, once subscribed to an observable, receives notification of new events emitted by the observable.

Operator: Operators operate on the data stream by transforming incoming data. The operator enables the developer to manipulate the data streams in a declarative manner.

Schedulers: The scheduler provides a mechanism to control the execution context of the observable.

After these core principles and key concepts of Reactive programming we will now list some implementation frameworks and libraries that support reactive programming.

RxJava/RxJS: RxJava and RxJS are reactive extensions for the Java and JavaScript programming languages, respectively. They provide a rich set of operators and utilities for composing asynchronous data streams and handling events.

FrTime: The frtime language supports declarative construction of reactive systems in a syntax very similar to that of Racket. It extends racket.

Akka Streams: Akka Streams is a toolkit for building stream processing applications using the Akka actor model. It offers a high-level DSL for defining complex stream processing topologies and supports backpressure handling out of the box.

ReactiveX: ReactiveX is a cross-platform library for reactive programming that supports multiple programming languages, including Java, JavaScript, C#, and Python. It provides a unified API for working with observables and operators across different platforms.

The above listed frameworks and libraries all have the downside of exposing one or two of the following problems that can arise in a typical two layer setup where the reactive system is build on top of an imperative language. Namely the *Reactive Thread Hijacking Problem* and the *Reactive/Imperative Impedance Mismatch*.(Vonder et al., 2020) After a short definition of both problems we will propose a novel reactive programming language, Haai. Haai is a reactive programming language that by design (one layer system) excludes the following two problems.

Reactive Thread Hijacking Problem

This happens in a two-layered reactive programming language when the reactive runtime or top layer calls a function in the bottom layer and that call never returns. The host has hijacked the control and the reactive runtime will not be able to regain control.

Reactive/Imperative Impedance Mismatch

When the reactive runtime in a two-layered reactive program used different side effects by calling functions on the first layer. It might be so that the order in which they are executed is not as expected. Creating faults in the programs that might not be easy to detect.

Levels of reactivity

What is the upper bound for long lasting computations in reactive programming? There is a trade off where more strick levels of reactivity allow for a less wide pallet of reactive programs to be expressed. For Haai the level of reactivity will be related to the computational time of the native reactors. In this implementation of the Haai virtual machine we only have native reactors that have a computational complexity of $O(1)$. This guarantees a strong level of reactivity as classified in this earlier work. (Vonder et al., 2020)

2.1 Haai

We are making use of the novel purely reactive language Haai. Haai is reactive all-the-way-trough. As such the only construct available for expressing computations is a reactor. A reactor

has sources or inputs and sinks or outputs. For each time the sources are updated the reactor will react and produce a new sink value. The number of sources and sinks are variable per reactor definition. In Haai there are no functions, only reactors. This allows us to build reactive programs without functions. (Oeyen et al., 2024)

2.2 Haai virtual machine

To make it possible for Haai code to work on multiple platforms one can make a virtual machine for each required platform. After the example of Remus, "a virtual machine that has been carefully designed to be usable for running reactive programs in low-powered computing environments." (Oeyen et al., 2022b) We implement in this thesis a virtual machine for Haai. The purpose of the Haai virtual machine (Hvm) is to enable the possibility of distributing Haai programs in a cluster of machines. There are multiple reasons to distribute a reactive program like: scalability, low latency and many others. Firstly the Hvm is a prototype to investigate distribution in general and to see if at all the language works well in a distributed setting. The Hvm is developed in the Elixir computer language. The chapter on distribution arguments our choice for Elixir. The Haai bytecode as compiled for Remus will be the representation of Haai reactors in the Hvm.

2.3 Remus instruction set

The Remus instruction set is the set of low-level operations that the virtual machine can execute. A Haai program or Reactor will after compilation exist out of a combination of instructions. Instructions form the Remus instruction set as listed here under. The Haai virtual machine is an implementation of this instruction set in Elixir.

[**I-ALLOCMONO**, **reactor**] Allocate memory for the given reactor.

[**I-LOOKUP**, **signal**] Lookup the value for signal in a signal table.

[**I-SUPPLY**, **from**, **destination**, **index**] Supply the value found in from to destination at index.

[**I-REACT**, **memory_location**] Apply the reaction required at memory_location.

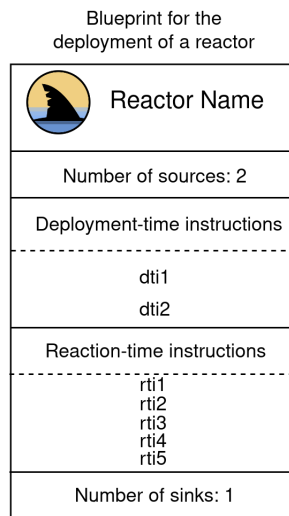
[**I-CONSUME**, **memory_location**, **index**] Move the value form REACT to the run-time-memory.

[**I-SINK**, **rti_location**, **sink_index**] Remove value from run-time-memory and store in the reactor sink.

2.4 Reactors

In Haai, a purely reactive computer language, the only construct available for expressing computations is a reactor. We can write a single reactor or combine multiple reactors in one 'main' reactor or program. The main program is build out of reactors, these can be native or user defined. The Haai language is by design resilient to the two problems mentioned earlier at the beginning of this chapter. No lifting of functions, only reactors. Every computation in a reactor is done by a reactor, some are native to the Hvm. The Hvm provides for example native reactors for basic arithmetic. We can represent a reactor as follows:

Figure 2.1: Graphical representation of a reactor



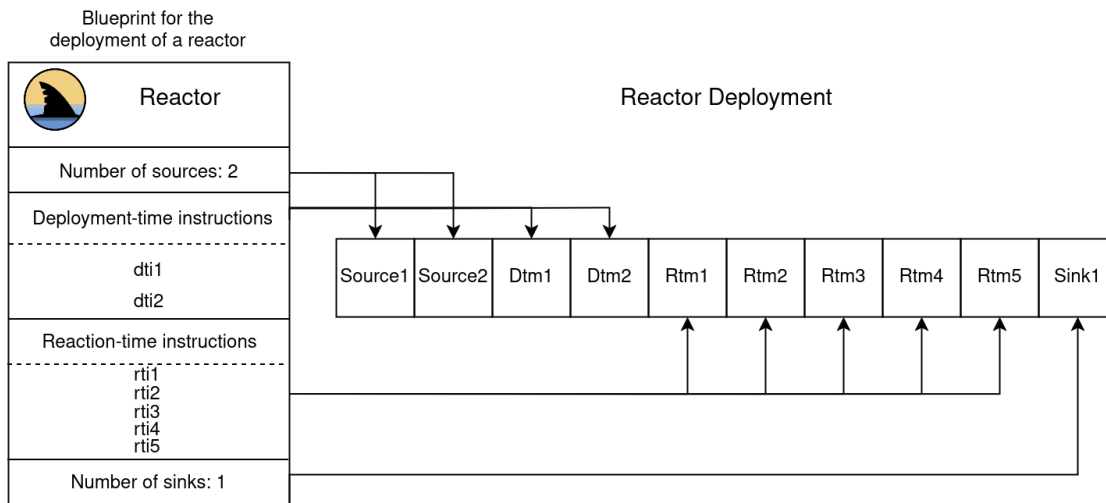
The deployment of a reactor reacts on a data stream, that data altered by the reaction is found in the sink of the reactor. when deploying a reactor we request a memory block fixed in size. The size of this memory block, the total memory required for the deployment of a reactor and is mainly influenced by the number of reactors a program exists of and the size of the run time instructions of these reactors.

2.5 Deploying a reactor

A reactor can be seen as a blue print for its deployment, similar how an object created from a class in object-oriented programming is an instantiation of the blueprint the class represents. The deployment of the reactor will request a certain size of memory. That total block of memory exists out of four parts: Sources, deployment time memory (dtm), reaction time memory (rtm) and the sinks. When we go from a reactor to a deployment we can find the following information in the reactor or blue print for the deployment. The number of sources and sinks, the deployment time instructions (dti) and the reaction time instructions (rti). For each source and each sink a position is reserved in the memory block we are constructing. From the deployment time instructions in the reactor we can prepare the deployment time memory (dtm). The Dtm part of the memoryblock we are constructing will reserve space for the deployment of the reactors found in the dti.

Now we have reserved space for the sources the sinks and dtm, the last step is to reserve space for the reaction time memory (rtm). The space required for rtm is equal to the number of instructions found in the reaction time instructions. For each instruction in rti there is one block reserved in rtm. A deployed reactor uses the four memory parts: sources, dtm, rtm and sinks to react according to the reaction time instructions found in rti. The following figure expresses the deployment of a reactor.

Figure 2.2: Reactor deployment, from blueprint to memory block



2.6 Running the reactor

Running the reactor encompasses starting a Haai virtual machine with all requested arguments for the Hvm to function. The Haai virtual machine needs the following arguments to start:

Reactor bytecode Bytecode obtained from the Haai compiler that compiles reactor definitions in the Haai language to a bytecode.

Connect source A function that the Hvm can use to connect some input to the sources of the reactor. At each start of the next iteration this function will be called to update the sources of the deployed reactor. For example connecting to a WebSocket server to receive the data it broadcasts.

handle sink The function that after each iteration is called to handle the new sinks produced by the deployed reactor. This function is tailored for the specific reactor that has been send as bytecode.

The following line demonstrates the call to the Elixir function `startVM` to start the Hvm and deploy the reactor found in `reactor_bytecode`. Three functions are passed as arguments, two for each source and one for the sinks.

```
startVM(node, reactor_bytecode, &connect_source1/0, &connect_source2/0, &handle_sinks/2)
```

2.7 Managing the reactors state

Managing the state of a reactor is done by one actor per program or reactor deployment. The actor is build around the `GenServer` module in Elixir which is a behavior module that provides a generic server implementation. In this implementation the name of the `GenServer` is `Memory`. `Memory` is a single actor that encapsulates multiple functions to respond to various messages or calls. Together these define the behaviour of the `Memory GenServer`.

Memory will mainly accept calls from the runtime instructions found in the reactor deployment. Available calls are mapped on the Haai instruction set. In addition there are calls to initialize or to get or set sources and sinks into or from the Memory GenServer. One instance of Memory represents the memory block of a deployed reactor.

The GenServer module in Elixir works with calls or casts, we use calls since they are synchronous requests to the GenServer. By calling Memory in a synchronous way we make sure that the values in Memory are only accessed by one operation at a time, assuring correct values.

2.8 Native reactors

The Haai virtual machine makes use of native reactors to perform simple arithmetic operations like addition, division, multiplication and subtraction. Native reactors are recognized by the Hvm and applied when required.

Haai is a one layer reactive programming language, contrary to two layer system where the top reactive layer can lift functions from the bottom imperative layer, Haai only uses reactors for all computation. The Hvm has a collection of frequent reactors. Thanks to the native reactors we do not need to implement the most basic reactors when composing a program. By design Haai prevents the Reactive/imperative impedance mismatch as explained in the chapter on Haai. Having the native reactors smartly covers the lack of lifting possibilities.

2.9 Special reactors

Specialized reactors are external to the reactive program. The specialized reactors allow us to perform operations with side effects, side effects are not allowed for a regular reactor or reactive program. In this implementation the special reactors are represented as Elixir functions one for each source and one for all sinks as explained in the Running the reactor section.

2.10 Example bytecode program

After parsing the bytecode that was compiled from the original Haai code with the Haai compiler into an Elixir readable format of nested lists a program could look like this. In this example we see two user defined reactors named `consonance` and `note.length` followed by the main reactor that represents the actual program that will be deployed.

```
co_nl = [
  [:consonance, 1, 1,
  [
    ["I-ALLOCMONO", :multiply]
  ],
  [
    ["I-LOOKUP", :ci],
    ["I-SUPPLY", ["%RREF", 1], ["%DREF", 1], 1],
    ["I-SUPPLY", ["%SRC", 1], ["%DREF", 1], 2],
    ["I-REACT", ["%DREF", 1]],
    ["I-CONSUME", ["%DREF", 1], 1],
    ["I-SINK", ["%RREF", 5], 1]]
  ],
```

```

[:note_length, 1, 1,
[
  ["I-ALLOCMONO", :divide], ["I-ALLOCMONO", :multiply]
],
[
  ["I-SUPPLY", 60000, ["%DREF", 1], 1],
  ["I-SUPPLY", ["%SRC", 1], ["%DREF", 1], 2],
  ["I-REACT", ["%DREF", 1]],
  ["I-LOOKUP", :lm],
  ["I-SUPPLY", ["%RREF", 4], ["%DREF", 2], 1],
  ["I-CONSUME", ["%DREF", 1], 1],
  ["I-SUPPLY", ["%RREF", 6], ["%DREF", 2], 2],
  ["I-REACT", ["%DREF", 2]],
  ["I-CONSUME", ["%DREF", 2], 1],
  ["I-SINK", ["%RREF", 9], 1]
],

[:main, 2, 2,
[
  ["I-ALLOCMONO", :consonance], ["I-ALLOCMONO", :note_length]
],
[
  ["I-SUPPLY", ["%SRC", 1], ["%DREF", 1], 1],
  ["I-REACT", ["%DREF", 1]],
  ["I-SUPPLY", ["%SRC", 2], ["%DREF", 2], 1],
  ["I-REACT", ["%DREF", 2]],
  ["I-CONSUME", ["%DREF", 1], 1],
  ["I-SINK", ["%RREF", 5], 1],
  ["I-CONSUME", ["%DREF", 2], 1],
  ["I-SINK", ["%RREF", 7], 2]
]
]
]

```

The above is an example for reactor.bytecode that is send into the Haai virtual machine. It is a compilation of the following Haai program defining two user defined reactors (consonance and note_length) and making use of two native reactors (:multiply (* in Haai) and :divide (/ in Haai)) and joining them together within the main reactor. The main reactor could be seen as the 'main' function in C. The starting point or the call that kicks off the programs execution.

```

(
  (defr (consonance f)
    (out (* ci f)))

    (defr (note_length bpm)
      (def q (/ 60000 bpm))
      (out (* lm q)))

    (defr (main f bpm)
      (out (consonance f))

```

```
(out (note_length bpm)))  
)
```


Chapter 3

Distribution

3.1 Shared-nothing architecture

To organize the distribution of the Haai virtual machine we have opted for a shared-nothing architecture. (Stonebraker, 1986) This allows us to build the cluster with the following characteristics, suitable for the first time distribution of the reactive language Haai.

Independent Nodes Each node can execute its own specific code and perform computations independently since no memory, storage or processing power is shared.

No shared state Each node manages its state independently. Cluster communication between nodes is done with messages. we use Erlang distributed mechanisms in the implementation.

Horizontal scalability Shared-nothing architectures are highly scalable horizontally, this makes it straightforward to add or remove nodes from the cluster. It allows us to run clusters of all sizes. Practically the cluster size will be bound to the available hardware.

Fault isolation Faults or failures in one node do not affect the operation of other nodes in the system. We can monitor node failure with Erlang distributed mechanisms and act accordingly.

Flexibility and Modularity Changes to one node do not impact the operation of other nodes, allowing for maintenance, upgrades, and system evolution of the cluster.‘

3.2 Erlang virtual machine

We implemented the shared-nothing architecture in Elixir a functional, concurrent, and fault-tolerant programming language built on the Erlang Virtual Machine (BEAM). The Erlang Virtual Machine provides a robust platform due to several key features and design principles:

Concurrency model Erlang has a lightweight concurrency model which is based on actor-like processes. Since each process runs independently and can communicate with other processes through message passing the Erlang virtual machine is inherently decentralized by design.

Fault tolerance The BEAM VM provides mechanisms for process supervision, fault isolation, and hot code swapping, allowing systems to recover from failures without affecting the overall operation of the cluster.

Distribution transparency Erlang’s distribution features enable processes to communicate effortlessly across network boundaries. Within an Erlang cluster, nodes can seamlessly interact through remote procedure calls (RPCs) and message passing. Different types of distribution models can be build on the Erlang platform.

Scalability Horizontal scaling or adding nodes to a cluster is straightforward. This makes it easy to scale clusters to need.

Tooling Common tasks are simplified thanks to a comprehensive tool set. tools like the Mix build tool and the OTP framework are core tools used in the implementation of the distributed Haai cluster.

The example program is explained in more detail in the use case chapter. The same program or reactor can be run many times. The distributed setup we make exists out of a number of elixir nodes that all send there messages to one sound server. The bigger the distributed network is the more messages that are send to the sound server. The sound server it self has virtually no limit to receive messages. The main parts of the distribution are taken care of by the underlying Beam virtual machine. One main reason to develop the Hvm on this platform is to be able to distribute in a solid and proven way.

3.3 Actor model

The actor system in Erlang implements the processes actor model. (De Koster et al., 2016) These actors are modeled as processes and can define there receive primitive to specify what messages it responds to. Thanks to the isolated turn principle the actor model guarantees to be free of data races and deadlocks by design. In Erlang actors are typically used when you need to encapsulate state and behavior within a single unit of computation., actors are spawned using constructs like GenServer, Task, or Agent within a single node. All actors communicate with each other using message passing within the same node. One can use these actors for modeling concurrent and stateful components in an application.

3.4 Distributed Erlang mechanisms

Distributed Erlang mechanisms come into play when you need to communicate between actors residing on different nodes within a cluster. Nodes communicate with each other using distributed Erlang protocols, allowing actors on different nodes to exchange messages. Distributed Erlang mechanisms provide fault tolerance, distribution transparency, and location transparency across the cluster.


Distributed Erlang provides location transparency, allowing actors to communicate regardless of their physical location within the cluster. This transparency hides the underlying network details, promoting loose coupling (Carreton et al., 2010) between nodes and actors. Actors on different nodes can interact seamlessly through message passing, without needing to be aware of the specific nodes they are communicating with. Loose coupling in distributed Erlang enables the system to scale horizontally by adding or removing nodes without significantly impacting the overall architecture.

3.5 Distributing the Hvm

The Elixir cluster exist out of a number of Elixir nodes. An idle Elixir nodes is not hungry for system resources. One can easily run a ten fold of nodes on one machine to simulate a cluster. Once a number of nodes are running one of the nodes will be selected as the controller or master node. The master node will be made aware of all nodes in the cluster. With all nodes connected to the master node but not to each other we can start the deployment of reactors. Deploying a reactor on the cluster comes down to starting the Hvm on that node with the bytecode of that requested reactor. This means that all nodes can run the same or a different reactor depending on the bytecode it received.

Once the cluster, one master node connected to all other nodes exists, the deployment can be summarized as follows, chose a node in the cluster, a reactor bytecode, functions to connect the sources and handle the sinks and fire of. The reactor or our reactive program is started and will keep on reacting for as long as the node exist in this cluster.

Figure 3.1: Distributing reactor deployment on Elixir cluster



Chapter 4

Evaluation

The use case of this thesis makes sound. We evaluate the distribution of the Haai reactors with sound. Each reactor is at given durations computing the next frequency to be sound on the system. This setup is very flexible as we can easily add or remove nodes from a cluster.

4.1 Distributed reactive melody generator

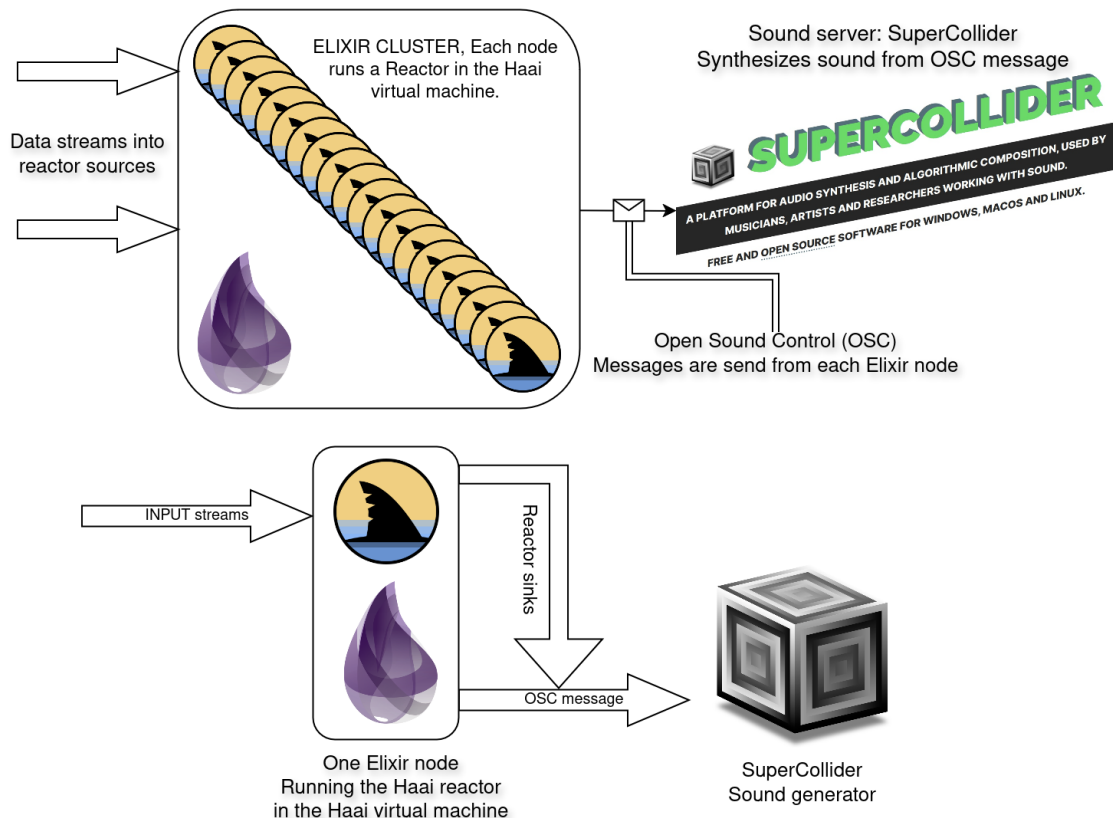
The distributed Haai cluster, existing out of many distributed reactors provides the sink values from the reactors as values inserted into OSC messages (Schmeder et al., 2010) that are send to a sound server to trigger the synthesis of some sound. It is the reactors that, by reacting on incoming streams of data, produce the values that result in a rhythmical flow of OSC messages to the sound server.

The flow of rhythmical messages is related to the duration of each sound. That same duration that makes the sound disappear after it was started determines the iteration speed for the reactor or program that runs on each individual node. Each reactor iterates over and over but each iteration has a defined duration, this duration is virtually equal to the duration for a sound calculated by the reactor. A reactor will calculate the new values, sleep for the calculated duration and then provide the new values in the sink. At this point the Hvm calls a function that it was given to handle the sink values. For this use case that function constructs the message and sends it to the sound server that propagates it to the requested synthesizer.

4.2 Sound server

The synthesizer that produces the sound is a Synthdef (the definition of a synthesizer) inside the sound server named SuperCollider. Supercollider is an open-source platform for audio synthesis and algorithmic composition. Developed by James McCartney in the late 1990s (Wilson et al., 2011), it has since become a powerful tool in the fields of music technology, computer music, and sound art. Supercollider provides a flexible and expressive environment for creating and manipulating sound in real-time, making it an invaluable resource for both research and artistic exploration. Some main features of super collider include:

Audio Synthesis Supercollider offers a wide range of synthesis techniques, including additive, subtractive, granular, and physical modeling synthesis. Users can create complex sounds by combining these techniques and modulating parameters in real-time.



Real-time Processing One of Supercollider’s key strengths is its ability to process audio in real-time. This makes it suitable for live performances, interactive installations, and other time-sensitive applications.

Algorithmic Composition Supercollider provides tools for generating music algorithmically, allowing users to create compositions based on mathematical algorithms, rulesets, or generative processes.

Integration with External Hardware Supercollider can interface with external MIDI controllers, audio interfaces, and other hardware devices, enabling users to incorporate physical instruments and sensors into their sound projects.

Community and Documentation Supercollider has a vibrant online community of users who share code, tutorials, and resources. Additionally, comprehensive documentation is available, including tutorials, reference guides, and examples to help users learn and master the software.

Supercollider has been and still is actively used in academia for many purposes, for example:

Research in Music Technology Digital signal processing, human-computer interaction, and machine learning for music. Researchers leverage its flexibility and programmability to prototype new algorithms, experiment with novel synthesis techniques, and investigate the perceptual and cognitive aspects of sound.

Composition and Sound Design Supercollider is used to create innovative works that push the boundaries of traditional music and sound art. Its ability to generate complex and evolving textures, as well as its support for algorithmic composition, makes it a valuable tool for exploring new sonic territories.

Teaching and Learning Supercollider is increasingly incorporated into music technology and computer music curricula at universities and colleges worldwide. It provides students with hands-on experience in sound synthesis, programming, and digital audio processing.

4.3 Sound server communication

The communication between the nodes in the Elixir cluster and the sound server is done over the network using the UDP network protocol. The sound server can handle a virtual unlimited amount of messages, but is practically bounded to the udp buffer size on the actual machine that runs the sound server. As such the Elixir cluster of reactors can be of any feasible size and the sound server will respond on any number of messages by synthesizing the requested sound.

4.4 Musical values

The reactor or program that runs on all nodes in the cluster does for each note the following. It calculates two values given two input streams of numbers. The first source is a stream of base frequencies, represented as a float number, the second source is again a stream of numbers representing a tempo in beats per minute (bpm)

Consonant notes

Consonant notes sound harmonious when played together. consonance can be expressed as a ratio between two notes or frequencies. Ratios like the perfect fifth (frequency ratio of 3:2) and the major third (frequency ratio of 5:4) are very commonly used ratios in a musical setting. One of the two reactors in the main program used for this use case calculates such a note for a given base note, both are expressed in there actual frequency represented as a float number.

Tempo

We can express the speed of a musical note progression in beats per minute (bpm). From that number we can then calculate the duration of individual notes in the musical performance.

We calculate the time in milliseconds. Since one minute is 60000 milliseconds one can calculate the length of a quarter note with $q = \frac{60000}{bpm}$ with the assumption that we use a 4/4 time signature. From the length of the quarter note it is easy to find other durations. For example, the duration of a halve note is double that of a quarter note or the duration of a eight note is half the time of a quarter note.

One of the two reactors in the main program used for this use case calculates exactly that, the length for a quarter note and improvises a deviser or multiplier to produce some duration.

Conclusion

The drmg project worked well, suggesting that the actor model is a good choice to distribute Haai in a shared-nothing cluster. The number of nodes that simulate the cluster can easily be a tenfold. The reactors itself have a very light computational cost existing out of some basic arithmetic operations. One Elixir node would consumes about 90MB of memory and virtually no processor time when idle and about 3% when performing the calculation.

Further work

Further work can be organized to explore different distribution models on the same Elixir set of nodes. Where all nodes did react independently of each other, in future work communication between nodes can be explored and more different programs or reactors can be setup and linked together... Study glitch avoidance

Bibliography

- Bainomugisha, E., Carreton, A. L., Cutsem, T. V., Mostinckx, S., & Meuter, W. D. (2013). A survey on reactive programming. *ACM Comput. Surv.*, 45(4), 52:1–52:34. <https://doi.org/10.1145/2501654.2501666>
- Carreton, A. L., Mostinckx, S., Cutsem, T. V., & Meuter, W. D. (2010). Loosely-coupled distributed reactive programming in mobile ad hoc networks. In J. Vitek (Ed.), *Objects, models, components, patterns, 48th international conference, TOOLS 2010, málaga, spain, june 28 - july 2, 2010. proceedings* (pp. 41–60). Springer. https://doi.org/10.1007/978-3-642-13953-6_3
- De Koster, J., Van Cutsem, T., & De Meuter, W. (2016). 43 years of actors: A taxonomy of actor models and their key properties. *Proceedings of the 6th International Workshop on Programming Based on Actors, Agents, and Decentralized Control*, 31–40. <https://doi.org/10.1145/3001886.3001890>
- Juric, S. (2024). *Elixir in action*. Simon; Schuster.
- Oeyen, B., De Koster, J., & De Meuter, W. (2022a). Reactive programming on the bare metal: A formal model for a low-level reactive virtual machine. *Proceedings of the 9th ACM SIGPLAN International Workshop on Reactive and Event-Based Languages and Systems*, 50–62.
- Oeyen, B., De Koster, J., & De Meuter, W. (2022b). Reactive programming on the bare metal: A formal model for a low-level reactive virtual machine. *Proceedings of the 9th ACM SIGPLAN International Workshop on Reactive and Event-Based Languages and Systems*, 50–62. <https://doi.org/10.1145/3563837.3568342>
- Oeyen, B., De Koster, J., & De Meuter, W. (2024). Reactive programming without functions. *The Art, Science, and Engineering of Programming*, 8(3), 11. <https://doi.org/10.22152/programming-journal.org/2024/8/11>
- Schmeder, A., Freed, A., & Wessel, D. (2010). Best practices for open sound control. *Linux Audio Conference*, 10.
- Stonebraker, M. (1986). The case for shared nothing. *IEEE Database Eng. Bull.*, 9(1), 4–9. <http://sites.computer.org/debull/86MAR-CD.pdf>
- Vonder, S. V. d., Renaux, T., Oeyen, B., Koster, J. D., & Meuter, W. D. (2020). Tackling the awkward squad for reactive programming: The actor-reactor model [ISSN: 1868-8969]. In R. Hirschfeld & T. Pape (Eds.), *34th european conference on object-oriented programming (ECOOP 2020)* (19:1–19:29). Schloss Dagstuhl–Leibniz-Zentrum für Informatik. <https://doi.org/10.4230/LIPIcs.ECOOP.2020.19>
- Wilson, S., Cottle, D., & Collins, N. (2011). *The supercollider book*. The MIT Press.